

# Module 4 Homework

For this homework, you will need the following datasets:

- [Green Taxi dataset \(2019 and 2020\)](#)
- [Yellow Taxi dataset \(2019 and 2020\)](#)
- [For Hire Vehicle dataset \(2019\)](#)

web\_to\_gcs.py:

```
import os
import requests
from google.cloud import storage

"""
Pre-reqs:
1. `pip install pandas google-cloud-storage`
2. Set GOOGLE_APPLICATION_CREDENTIALS to your project/service-account key
3. Set GCP_GCS_BUCKET as your bucket or change default value of BUCKET
"""

# 初始化下载链接
init_url = 'https://github.com/DataTalksClub/nyc-tlc-data/releases/download/'
# 设置 GCS Bucket 名称
BUCKET = "database_sliu"

def upload_to_gcs(bucket_name, object_name, local_file):
    """
    上传文件到 GCS
    Ref: https://cloud.google.com/storage/docs/uploading-objects#storage-upload-object-python
    """
    CREDENTIALS_FILE = "gcs.json"
    client = storage.Client.from_service_account_json(CREDENTIALS_FILE)
    bucket = client.bucket(bucket_name)
    blob = bucket.blob(object_name)
    blob.upload_from_filename(local_file)
    print(f"Uploaded to GCS: gs://{bucket_name}/{object_name}")

def web_to_gcs(year, service):
    for i in range(12):
        month = f"{i+1:02d}" # 格式化月份, 确保两位数
        file_name = f"{service}_tripdata_{year}-{month}.csv.gz"
        request_url = f"{init_url}{service}/{file_name}"

        # 下载文件
        print(f"Downloading: {request_url}")
        r = requests.get(request_url)

        if r.status_code == 200:
            with open(file_name, 'wb') as f:
                f.write(r.content)
            print(f"Downloaded: {file_name}")
        else:
```

```

        print(f"Failed to download {file_name}, status code:
{r.status_code}")
        continue # 如果下载失败, 跳过当前月份

# 直接上传 .csv.gz 到 GCS
gcs_path = f"{service}/{file_name}"
upload_to_gcs(BUCKET, gcs_path, file_name)

# 删除本地文件
os.remove(file_name)
print(f"Deleted local file: {file_name}")

# 调用函数下载并上传到 GCS
# web_to_gcs('2019', 'green')
# web_to_gcs('2020', 'green')
web_to_gcs('2019', 'fhv')

```

load data sql in GCP:

```

-- LOAD DATA INTO `dbt-learn-452007.trips_data_all.fhv_tripdata`
-- FROM FILES(
--   format = 'CSV',
--   uris = ['gs://database_sliu/fhv/*.csv.gz']
-- );

select count(*) from `dbt-learn-452007.trips_data_all.fhv_tripdata`;

```

## Before you start

1. Make sure you, **at least**, have them in GCS with a External Table **OR** a Native Table - use whichever method you prefer to accomplish that (Workflow Orchestration with [pandas-gbq](#), [dlt for gcs](#), [dlt for BigQuery](#), [gsutil](#), etc)
2. You should have exactly `7,778,101` records in your Green Taxi table
3. You should have exactly `109,047,518` records in your Yellow Taxi table
4. You should have exactly `43,244,696` records in your FHV table
5. Build the staging models for green/yellow as shown in [here](#)
6. Build the dimension/fact for taxi\_trips joining with `dim_zones` as shown in [here](#)

**Note:** If you don't have access to GCP, you can spin up a local Postgres instance and ingest the datasets above

## Question 1: Understanding dbt model resolution

- database.schema.table
- myproject.raw\_nyc\_tripdata.ext\_green\_taxi

Provided you've got the following sources.yaml

#如果为设置，则设置为默认

```
version: 2

sources:
  - name: raw_nyc_tripdata
    database: "{{ env_var('DBT_BIGQUERY_PROJECT', 'dtc_zoomcamp_2025') }}"
    schema:   "{{ env_var('DBT_BIGQUERY_SOURCE_DATASET', 'raw_nyc_tripdata') }}"
    tables:
      - name: ext_green_taxi
      - name: ext_yellow_taxi
```

with the following env variables setup where `dbt` runs:

```
export DBT_BIGQUERY_PROJECT=myproject
export DBT_BIGQUERY_DATASET=my_nyc_tripdata
```

What does this .sql model compile to?

```
select *
from {{ source('raw_nyc_tripdata', 'ext_green_taxi') }}
```

- `select * from dtc_zoomcamp_2025.raw_nyc_tripdata.ext_green_taxi`
- `select * from dtc_zoomcamp_2025.my_nyc_tripdata.ext_green_taxi`
- `select * from myproject.raw_nyc_tripdata.ext_green_taxi`
- `select * from myproject.my_nyc_tripdata.ext_green_taxi`
- `select * from dtc_zoomcamp_2025.raw_nyc_tripdata.green_taxi`

## Question 2: dbt Variables & Dynamic Models

Say you have to modify the following `dbt_model ( fct_recent_taxi_trips.sql )` to enable Analytics Engineers to dynamically control the date range.

- In development, you want to process only **the last 7 days of trips**
- In production, you need to process **the last 30 days** for analytics

```
select *
from {{ ref('fact_taxi_trips') }}
where pickup_datetime >= CURRENT_DATE - INTERVAL '30' DAY
```

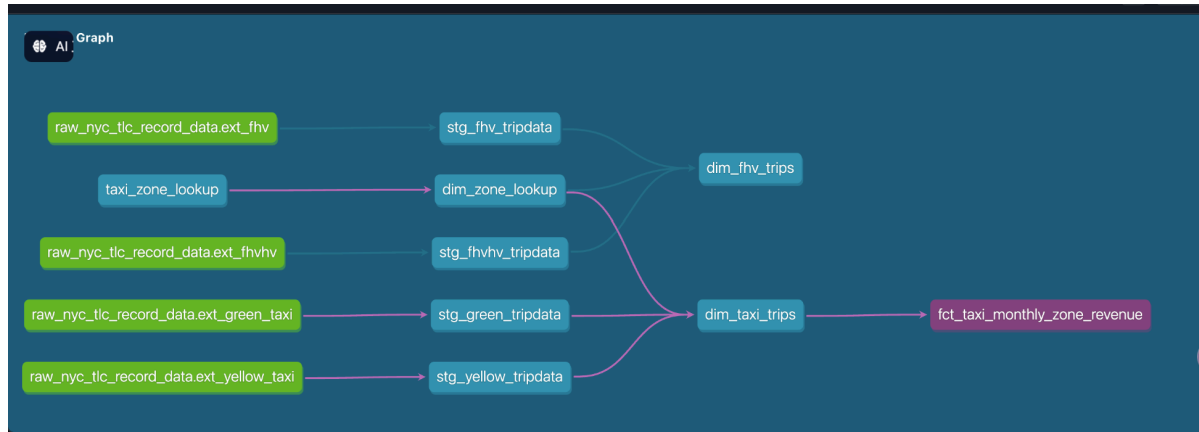
What would you change to accomplish that in a such way that command line arguments takes precedence over ENV\_VARS, which takes precedence over DEFAULT value?

- Add `ORDER BY pickup_datetime DESC` and `LIMIT {{ var("days_back", 30) }}`
- Update the WHERE clause to `pickup_datetime >= CURRENT_DATE - INTERVAL '{{ var("days_back", 30) }}' DAY`

- Update the WHERE clause to `pickup_datetime >= CURRENT_DATE - INTERVAL '{{ env_var("DAYS_BACK", "30") }}' DAY`
- **Update the WHERE clause to** `pickup_datetime >= CURRENT_DATE - INTERVAL '{{ var("days_back", env_var("DAYS_BACK", "30")) }}' DAY`
- Update the WHERE clause to `pickup_datetime >= CURRENT_DATE - INTERVAL '{{ env_var("DAYS_BACK", var("days_back", "30")) }}' DAY`

### Question 3: dbt Data Lineage and Execution

Considering the data lineage below **and** that `taxi_zone_lookup` is the **only** materialization build (from a .csv seed file):



Select the option that does **NOT** apply for materializing `fct_taxi_monthly_zone_revenue`:

- `dbt run`
- `dbt run --select +models/core/dim_taxi_trips.sql+ --target prod`
- `dbt run --select +models/core/fct_taxi_monthly_zone_revenue.sql`
- `dbt run --select +models/core/`
- `dbt run --select models/staging/+`

### Question 4: dbt Macros and Jinja

Consider you're dealing with sensitive data (e.g.: [PII](#)), that is **only available to your team and very selected few individuals**, in the `raw` layer of your DWH (e.g: a specific BigQuery dataset or PostgreSQL schema),

- Among other things, you decide to obfuscate/masquerade that data through your staging models, and make it available in a different schema (a `staging` layer) for other Data/Analytics Engineers to explore
- And **optionally**, yet another layer (`service` layer), where you'll build your dimension (`dim_`) and fact (`fct_`) tables (assuming the [Star Schema dimensional modeling](#)) for Dashboarding and for Tech Product Owners/Managers

You decide to make a macro to wrap a logic around it:

```
{% macro resolve_schema_for(model_type) -%}

    {% set target_env_var = 'DBT_BIGQUERY_TARGET_DATASET' -%}
    {% set staging_env_var = 'DBT_BIGQUERY_STAGING_DATASET' -%}

    {% if model_type == 'core' -%} {{ env_var(target_env_var) -}}
    {% else -%}                      {{ env_var(staging_env_var,
env_var(target_env_var)) -}}
    {% endif -%}

{% endmacro %}
```

And use on your staging, dim\_ and fact\_ models as:

```
{{ config(
    schema=resolve_schema_for('core'),
) }}
```

That all being said, regarding macro above, **select all statements that are true to the models using it:**(I select the false one)

- Setting a value for `DBT_BIGQUERY_TARGET_DATASET` env var is mandatory, or it'll fail to compile
- **Setting a value for `DBT_BIGQUERY_STAGING_DATASET` env var is mandatory, or it'll fail to compile**
- When using `core`, it materializes in the dataset defined in `DBT_BIGQUERY_TARGET_DATASET`
- When using `stg`, it materializes in the dataset defined in `DBT_BIGQUERY_STAGING_DATASET`, or defaults to `DBT_BIGQUERY_TARGET_DATASET`
- When using `staging`, it materializes in the dataset defined in `DBT_BIGQUERY_STAGING_DATASET`, or defaults to `DBT_BIGQUERY_TARGET_DATASET`

## Serious SQL

Alright, in module 1, you had a SQL refresher, so now let's build on top of that with some serious SQL.

These are not meant to be easy - but they'll boost your SQL and Analytics skills to the next level. So, without any further do, let's get started...

You might want to add some new dimensions `year` (e.g.: 2019, 2020), `quarter` (1, 2, 3, 4), `year_quarter` (e.g.: 2019/Q1, 2019-Q2), and `month` (e.g.: 1, 2, ..., 12), **extracted from pickup\_datetime**, to your `fact_taxi_trips` OR `dim_taxi_trips.sql` models to facilitate filtering your queries

## Question 5: Taxi Quarterly Revenue Growth

1. Create a new model `fact_taxi_trips_quarterly_revenue.sql`
  2. Compute the Quarterly Revenues for each year for based on `total_amount`
  3. Compute the Quarterly YoY (Year-over-Year) revenue growth
- e.g.: In 2020/Q1, Green Taxi had -12.34% revenue growth compared to 2019/Q1

- e.g.: In 2020/Q4, Yellow Taxi had +34.56% revenue growth compared to 2019/Q4

Considering the YoY (Year over Year) Growth in 2020, which were the yearly quarters with the best (or less worse) and worst results for green, and yellow

- green: {best: 2020/Q2, worst: 2020/Q1}, yellow: {best: 2020/Q2, worst: 2020/Q1}
- green: {best: 2020/Q2, worst: 2020/Q1}, yellow: {best: 2020/Q3, worst: 2020/Q4}
- green: {best: 2020/Q1, worst: 2020/Q2}, yellow: {best: 2020/Q2, worst: 2020/Q1}
- **green: {best: 2020/Q1, worst: 2020/Q2}, yellow: {best: 2020/Q1, worst: 2020/Q2}**
- green: {best: 2020/Q1, worst: 2020/Q2}, yellow: {best: 2020/Q3, worst: 2020/Q4}

```
-- select service_type, year_quarter, sum(total_amount)
-- from `dbt_sliu.fact_trips`
-- where extract(year from pickup_datetime) in (2019, 2020)
-- group by service_type, year_quarter;
```

## Question 6: P97/P95/P90 Taxi Monthly Fare

JOB INFORMATION		RESULTS	CHART	JSON	EXECUTION DETAILS		EXECUTION GRAPH	
Row	service_type	year	month	p97	p95	p90		
1	Green	2020	4	55	45	26.5		
2	Yellow	2020	4	32	25.5	19		

```
with filtered_trips as (
  select
    service_type,
    extract(year from pickup_datetime) as year,
    extract(month from pickup_datetime) as month,
    fare_amount
  from `dbt_sliu.fact_trips`
  where fare_amount > 0
    and trip_distance > 0
    and payment_type_description in ('Cash', 'Credit card')
),
percentiles as (
  select
    service_type,
    year,
    month,
    approx_quantiles(fare_amount, 100)[SAFE_OFFSET(97)] as p97,
    approx_quantiles(fare_amount, 100)[SAFE_OFFSET(95)] as p95,
    approx_quantiles(fare_amount, 100)[SAFE_OFFSET(90)] as p90
  from filtered_trips
  group by service_type, year, month
)
select *
from percentiles
where year = 2020 and month = 4;
```

1. Create a new model `fct_taxi_trips_monthly_fare_p95.sql`
2. Filter out invalid entries (`fare_amount > 0`, `trip_distance > 0`, and `payment_type_description in ('Cash', 'Credit Card')`)
3. Compute the **continuous percentile** of `fare_amount` partitioning by `service_type`, year and month

Now, what are the values of `p97`, `p95`, `p90` for Green Taxi and Yellow Taxi, in April 2020?

- green: {p97: 55.0, p95: 45.0, p90: 26.5}, yellow: {p97: 52.0, p95: 37.0, p90: 25.5}
- **green: {p97: 55.0, p95: 45.0, p90: 26.5}, yellow: {p97: 31.5, p95: 25.5, p90: 19.0}**
- green: {p97: 40.0, p95: 33.0, p90: 24.5}, yellow: {p97: 52.0, p95: 37.0, p90: 25.5}
- green: {p97: 40.0, p95: 33.0, p90: 24.5}, yellow: {p97: 31.5, p95: 25.5, p90: 19.0}
- green: {p97: 55.0, p95: 45.0, p90: 26.5}, yellow: {p97: 52.0, p95: 25.5, p90: 19.0}

## Question 7: Top #Nth longest P90 travel time Location for FHV

Prerequisites:

- Create a staging model for FHV Data (2019), and **DO NOT** add a deduplication step, just filter out the entries where `where dispatching_base_num is not null`
- Create a core model for FHV Data (`dim_fhv_trips.sql`) joining with `dim_zones`. Similar to what has been done [here](#)
- Add some new dimensions `year` (e.g.: 2019) and `month` (e.g.: 1, 2, ..., 12), based on `pickup_datetime`, to the core model to facilitate filtering for your queries

Now...

1. Create a new model `fct_fhv_monthly_zone_traveltime_p90.sql`
2. For each record in `dim_fhv_trips.sql`, compute the `timestamp_diff` in seconds between `dropoff_datetime` and `pickup_datetime` - we'll call it `trip_duration` for this exercise
3. Compute the **continuous** `p90` of `trip_duration` partitioning by year, month, `pickup_location_id`, and `dropoff_location_id`

For the Trips that **respectively** started from `Newark Airport`, `SoHo`, and `Yorkville East`, in November 2019, what are **dropoff\_zones** with the 2nd longest `p90` `trip_duration` ?

- **LaGuardia Airport, Chinatown, Garment District**
- LaGuardia Airport, Park Slope, Clinton East
- LaGuardia Airport, Saint Albans, Howard Beach
- LaGuardia Airport, Rosedale, Bath Beach
- LaGuardia Airport, Yorkville East, Greenpoint

```
WITH trip_p90 AS (
  SELECT
    pickup_zone,
    dropoff_zone,
    APPROX_QUANTILES(trip_duration, 100)[OFFSET(90)] AS p90_trip_duration
  FROM `dbt_sliu.dim_fhv_trips`
)
```

```

WHERE year = 2019
      AND month = 11
      AND pickup_zone IN ('Newark Airport', 'SoHo', 'Yorkville East')
GROUP BY pickup_zone, dropoff_zone
),

ranked_trips AS (
  SELECT *,
         RANK() OVER (PARTITION BY pickup_zone ORDER BY p90_trip_duration DESC) AS
rnk
  FROM trip_p90
)

SELECT pickup_zone, dropoff_zone, p90_trip_duration
FROM ranked_trips
WHERE rnk = 2;

```

#### Query results

JOB INFORMATION		RESULTS	CHART	JSON	EXECUTION DETAILS	EXECUTION GRAPH
Row	pickup_zone ▼	dropoff_zone ▼	p90_trip_duration ▼			
1	Yorkville East	Garment District	13846			
2	Newark Airport	LaGuardia Airport	7251			
3	SoHo	Greenwich Village South	26928			

## Submitting the solutions

- Form for submitting: <https://courses.datatalks.club/de-zoomcamp-2025/homework/hw4>

## Solution

- To be published after deadline